

Didactic Manual

big.bi $\rightarrow$ bi.gbi

Articulatory Simulation of the Verbal Transformation

Reproduction of the demonstrations described in

“Why can big.bi be changed to bi.gbi?”

A mathematical model of syllabification and articulatory synthesis”

arXiv:2307.02299 (Berthommier, July 2023)

Model: VLAM (Maeda, 7 parameters)

Code: Python 3.9+ (MIT License)

Dependencies: numpy, scipy, matplotlib, opencv-python

Repository: bigbi-article-repo

This manual provides step-by-step instructions for reproducing the verbal transformation simulations, including the polar trajectory visualization, the U-shaped Tp variation, and the comparison with the article’s Figure 4.

Contents

1	Introduction	3
1.1	What this repository provides	3
1.2	The VLAM model	3
1.3	Article parameters	4
2	The Mathematical Model	5
2.1	Polar coordinates and targets	5
2.2	Arc trajectories	5
2.3	Anchoring vowels: Vo and Ve	5
2.4	Superposition	5
3	The Syllable Graph Architecture	6
3.1	CV graph	6
3.2	CVC graph	6
3.3	CCV graph (clusters)	6
3.4	The fusion mechanism	6
4	The δ_o and δ_e Parameters	7
4.1	Implementation	7
4.2	Effect on the anchoring vowels	7
4.3	The V_e plateau — revoked	7
5	Reproduction Guide	8
5.1	Installation	8
5.2	Demonstration 1: big.bi → bi.gbi (Article §4)	8
5.3	Demonstration 2: U-shaped T_p variation	8
5.4	Demonstration 3: Comparison with the article	9
6	Polar Trajectory Reconstruction	9
6.1	Reconstruction rules	9
6.2	Display curvature and ν convention	10
6.3	Comparison with the article’s Figure 4	10
7	Results and Duration Analysis	10
7.1	Duration reduction	10
7.2	Perceptual salience of the schwa @	11
8	Code Structure	11

8.1	Modified files (vs original timit-to-Maeda)	11
8.2	Simulation scripts	11
9	Quick Demo	11
10	Article Demonstrations	12
10.1	Demo 1: /ibia/ VCV synthesis (Figure 1)	12
10.2	Demo 2: Locus equations (Figure 3)	13
10.2.1	Root cause of the previous /b-/d/ inversion	13
10.3	Demo 4: Trough effect (Figure 1b)	14
10.4	Demo 5: Six consonant cluster pairs	15
10.5	Original article simulations (exact reproduction)	15
10.5.1	Figure 1: /ibia/ with T=100ms	16
10.5.2	Figure 4: big.bi and bi.gbi	16
10.5.3	Supplement: /ib/ → /bi/	16
11	Interactive HTML Simulations	16
11.1	big.bi → bi.gbi (T=16, starts at COEFCEN=0.5)	16
11.2	ib ib → bi bi (NON-REVERSIBLE U-shape)	17
11.3	ib ib → bi bi (U-shape, reversible)	17
12	Conclusion	17

1 Introduction

This manual documents the Python implementation of the articulatory speech synthesis model described in arXiv:2307.02299 [1]. The model, based on VLAM (a Maeda articulatory synthesizer with 7 parameters), simulates the **verbal transformation** from **big.bi** to **bi.gbi** — a phenomenon where the syllable boundary shifts as the pause between syllables is shortened and the anchoring vowels (V_o , V_e) approach the main vowel V .

The article’s central claim (§4) is:

“When $\delta_o = \delta_e$ tends to 1 the graph reorganization can occur because two successive superimposed segments of duration $2 \times 2T$ anchored to the same vowel $V_o = V_e = V$ can fuse into a single one of duration $3T$.”

This repository provides the code to reproduce this transformation step by step, with audio synthesis, polar trajectory visualization, and accessible video output.

1.1 What this repository provides

1. The `synthSYL` package (modified from [FBerthommier/timit-to-Maeda](#)) with the δ_o and δ_e parameters added to `panphon_pipeline()`.
2. The VLAM synthesizer (`vlam.py`) for audio generation from Maeda parameters.
3. Simulation scripts for three demonstrations:
 - **big.bi** → **bi.gbi** (the article’s main result, §4)
 - **ib ib** → **bi bi** (the classical transformation, §3)
 - U-shaped T_p variation for perceptual demonstration
4. Polar trajectory reconstruction (two-branch: vocalic z_v and consonantal z_c), using primitives from [FBerthommier/covtl-pipeline](#).
5. Dual-panel video generation (sagittal vocal tract | polar plot) with synchronized audio and phoneme labels.

1.2 The VLAM model

VLAM (Vocal tract Learner Articulatory Model) is the Maeda model with 7 articulatory parameters:

Index	Symbol	Description
0	J	Jaw opening
1	B	Tongue body position
2	D	Tongue dorsum height
3	T	Tongue tip position
4	LP	Lip protrusion/rounding
5	LH	Vertical lip opening
6	Hy	Larynx height

Indexing convention. Two numbering conventions coexist and must not be mixed:

- the **article** numbers the parameters 1–7 (1=Jaw, 2=Body, 3=Dorsum, 4=Tip, 5=LipP, 6=LipH, 7=Hy); its selection vectors S_c and S_v (e.g. $S_c = \{1, 2, 6\}$) use this 1-based numbering;
- the **code** (and the table above) numbers them 0–6; the registry selection vectors are 0-based, e.g. /b/ has `params=[0,1,5]`.

Correspondence: article $\{1, 2, 6\} = \text{code } \{0, 1, 5\} = \{\text{Jaw, Body, LipH}\}$; article $\{1, 2, 3, 4\} = \text{code } \{0, 1, 2, 3\} = \{\text{Jaw, Body, Dorsum, Tip}\}$; article $S_v = \{3, 4, 5, 7\} = \text{code } \{2, 3, 4, 6\} = \{\text{Dorsum, Tip, LipP, Hy}\}$. Throughout this manual, S_c/S_v sets are quoted in the article’s 1-based numbering (as in the article), while parameter indices mentioned in the text refer to the 0-based code indices unless stated otherwise.

The coordination function (article Eq. 1) maps a polar target (ρ, θ) to the 7 Maeda parameters:

$$P_i = \Omega_i + \rho \cdot \Psi_{1,i} \cdot \cos(\Psi_{2,i} - \theta) \quad (1)$$

where Ω , Ψ_1 , and Ψ_2 are calibration constants (article Table 1). The CO matrix in `synthSYL/constants.py` matches the article’s Table 1 almost exactly (only the Tip parameter c_0 differs slightly: -2.75 in `synthSYL` vs -3.0 in the article).

1.3 Article parameters

The article (§2.1) specifies the following arc parameters:

- $K_{\text{voy}} = 30$ — vowel arc curvature
- $K = 10$ — consonant arc curvature
- $P_{\text{exp}} = 1$ — velocity profile exponent ($\rho(t) = \cos(\theta(t)/2)$, not $\cos^2$)
- $T = 16$ — time step (didactic mode)
- $\nu = -1$ — coarticulation index

These are set in `synthSYL/constants.py`:

```
DEFAULT_T: int = 16
DEFAULT_K: int = 10           # consonant arc curvature
DEFAULT_KVOY: int = 30       # vowel arc curvature
DEFAULT_PEXP: int = 1        # velocity profile exponent
DEFAULT_NU: int = -1
```

The article notes: “When K is large ($K = 30$ for vowel arcs and $K = 10$ for consonant arcs), the trajectories become straight lines.” With $P_{\text{exp}} = 1$, the velocity profile follows $\cos(\theta/2)$ (a cosine half-bell) rather than $\cos^2(\theta/2)$ (a Hanning window).

Note on display curvature: The polar trajectory visualization uses $K = 30$ for both vocalic and consonantal branches (the `covtl`-pipeline reference display curvature), not the engine’s $K = 10$ / $K_{\text{voy}} = 30$. This is a deliberate visualization choice that produces moderate, visible curvature for both branches.

2 The Mathematical Model

2.1 Polar coordinates and targets

Each phoneme target is a point (ρ, θ) in the complex plane:

- **Vowels:** $\rho \in [0, 1]$, $\theta \in [0, 2\pi)$
- **Consonants:** $\rho \in [0, 1.2]$, $\theta \in [0, 2\pi)$

For example, the vowel /i/ is at $(\rho = 0.9, \theta = 5\pi/3)$ and the consonant /b/ is at $(\rho = 1.2, \theta = \pi/3)$.

2.2 Arc trajectories

Motion between two targets (ρ_1, θ_1) and (ρ_2, θ_2) is produced by the arc function (article Eq. 2):

$$z(t) = (1 - r_k) \cdot \rho_1 e^{i\theta_1} + r_k \cdot \rho_2 \cdot e^{i(\theta_2 + \frac{\nu}{K} \cdot \text{th})} \quad (2)$$

where $r_k = \cos^{P_{\text{exp}}}(\text{th}/2)$ is the velocity profile, $\text{th} \in [0, \pi]$ is a half-turn sweep, $\nu = \pm 1$ controls the curvature direction, and K is the curvature constant ($K = 30$ for display, $K = 10$ for consonant arcs, $K_{\text{voy}} = 30$ for vowel arcs in the synthesis engine).

Important: The article uses $P_{\text{exp}} = 1$, so $r_k = \cos(\text{th}/2)$ (a cosine half-bell, *without* the power of 2). This is the form used for all demonstrations in this repository. The original **synthSYL** package used $P_{\text{exp}} = 2$ ($r_k = \cos^2(\text{th}/2)$, a Hanning window), which has been corrected to match the article.

2.3 Anchoring vowels: Vo and Ve

The article (§2.2) introduces two anchoring vowels:

$$V_o = (\delta_o \cdot \rho_V, \theta_V) \quad (3)$$

$$V_e = (\delta_e \cdot \rho_V, \theta_V) \quad (4)$$

- V_o (onset): anticipation of the vowel before the consonant. The value of δ_o determines the degree of anticipation.
- V_e (end): centralization of the coda vowel (schwa-like vocoid). Modeled with $\delta_e = 0.5$ (salient) or $\delta_e = 0.7$ (audible).

When $\delta_o = \delta_e \rightarrow 1$, both V_o and V_e approach V (the full vowel), enabling the fusion of adjacent syllable segments.

2.4 Superposition

The final parameter frame is a blend (article Eq. 4):

$$P(t) - \Omega = \text{Re} [S_v \cdot \Psi \cdot \bar{z}_v(t) + S_c \cdot \Psi \cdot \bar{z}_c(t)] \quad (5)$$

where S_v and S_c are selection vectors (binary masks over the 7 parameters), $z_v(t)$ is the vocalic trajectory, and $z_c(t)$ is the consonantal trajectory.

3 The Syllable Graph Architecture

3.1 CV graph

A CV syllable (e.g., /bi/) has 3 nodes and 4 arcs:

- Nodes: $V = \{V_o, C, V\}$
- Locations: $L = \{(\delta_o \rho_V, \theta_V), (\rho_C, \theta_C), (\rho_V, \theta_V)\}$
- Arcs: $E = \{(V_o, C), (C, V), (V_e, V), (V, V)\}$
- Durations: $T_p = \{T_1, T_2, 2T_2, T\}$

3.2 CVC graph

A CVC syllable (e.g., /big/) has 5 nodes, symmetric with respect to the CV graph:

- Nodes: $V = \{V_o, C_1, V, C_2, V_e\}$
- Locations: $L = \{(\delta_o \rho_V, \theta_V), (\rho_{C_1}, \theta_{C_1}), (\rho_V, \theta_V), (\rho_{C_2}, \theta_{C_2}), (\delta_e \rho_V, \theta_V)\}$

3.3 CCV graph (clusters)

A CCV syllable (e.g., /gbi/) has 4 nodes and 5 arcs:

- Nodes: $V = \{V_o, C_1, C_2, V\}$
- Arcs: $E = \{(V_o, C_1), (C_1, C_2), (C_2, V), (V_e, V), (V, V)\}$
- The consonantal branch visits C_1 then C_2 in sequence.

3.4 The fusion mechanism

The article (§4) describes the fusion:

“When $\delta_o = \delta_e$ tends to 1 the graph reorganization can occur because two successive superimposed segments of duration $2 \times 2T$ anchored to the same vowel $V_o = V_e = V$ can fuse into a single one of duration $3T$.”

The gain is twofold:

1. **Consonant clustering:** reduction of one period T (from 2 separate consonant gestures to 1 fused cluster gesture).
2. **Node loss:** the V node between the two syllables disappears (since $V_o = V_e = V$), reducing the graph complexity.

Additionally, the selection vector switches (article’s 1-based numbering; code indices subtract 1):

- Before fusion: $S_c = \{1, 2, 3, 4\}$ for /g/ and $S_c = \{1, 2, 6\}$ for /b/ (two separate gestures)
- After fusion: $S_c = \{1, 2, 3, 6\}$ for /gb/ (one fused gesture)

4 The δ_o and δ_e Parameters

4.1 Implementation

The δ_o and δ_e parameters are added to the `panphon_pipeline()` function as optional keyword arguments:

```
from synthSYL import panphon_pipeline

result = panphon_pipeline(
    "big bi",
    T=16,
    short_pause_duration_factor=1.0,
    delta_o=0.5,    #  $V_o = (0.5 * \rho_V, \theta_V)$ 
    delta_e=0.5,    #  $V_e = (0.5 * \rho_V, \theta_V)$ 
)
```

The default values are $\delta_o = \delta_e = \text{COEFCEN} = 0.5$, preserving backward compatibility with the original package.

4.2 Effect on the anchoring vowels

Table 1 shows the effect of δ on the anchoring vowels for the vowel /i/ ($\rho_V = 0.9$, $\theta_V = 5\pi/3$):

Table 1: Effect of δ on V_o and V_e (for /i/)			
δ	$V_o = V_e \rho$	Perception	Article term
0.50	0.450	Salient schwa @	“salient”
0.70	0.630	Audible schwa	“audible”
0.85	0.765	Residual schwa	—
0.95	0.855	$V_o \approx V_e \approx V$	—
1.00	0.900	$V_o = V_e = V$ (fusion)	—

4.3 The V_e plateau — revoked

An earlier revision of this repository added a full-amplitude V_e plateau before the decay block (`trajectory.py`, `_handle_b_terminal()`), reasoning that “ V_e is a NODE” (§2.2) and should therefore be held like V . This addition has been **revoked** (author ruling 2026-10-07): a reduced word-end vowel is not pronounced as a tensed vowel — one does not pronounce “ibe” — so there is NO held @. The engine now behaves like the reference `timit-to-Maeda`: the trajectory *reaches* the V_e anchor and decays immediately, with no extra hold. The reference `_handle_b_terminal()` chains the decay directly:

```
elif A.kind == "V":
    if B.kind == "pause":
        pt_A = _safe_pt(A)
        _append_decay_block(blocks, block_info, pt_A, T_cons, pre_amp=1.0)
```

Our engine now matches this (the final vowel of a CV word is still held through the regular last-anchor hold; only the extra pre-decay V_e plateau is gone). All reference reproductions remain bit-identical (0.000e+00): /ibia/ never had a V_e plateau, and the words whose acoustics change (e.g. /ib/) are not part of the stored reference gates.

5 Reproduction Guide

5.1 Installation

```
# Clone the repository
git clone https://github.com/<your-org>/bigbi-article-repo.git
cd bigbi-article-repo

# Install Python dependencies
pip install -r requirements.txt

# (Optional) Install ffmpeg for video generation
sudo apt install ffmpeg # Linux
# brew install ffmpeg   # macOS
```

5.2 Demonstration 1: big.bi → bi.gbi (Article §4)

This is the main demonstration, showing the verbal transformation with $T = 16$ (didactic mode), starting at $\delta = 0.50$ for a salient boundary schwa @. **Dot form** (author ruling 2026-10-07): big.bi is synthesized as the dotted word **big.bi** — the ' between /g/ and /b/ is a C.C syllable boundary with *no pause* (the reference Timit-to-Maeda semantics: the Ve of “big” is coarticulated with the /i/ of “bi”, the boundary anchor weighted by COEFCEN). There is therefore *no* T_p /pause dimension in this sweep: δ alone drives the transformation (boundary anchor = $\delta \cdot \rho_i$; at $\delta = 1$ the anchor is the full /i/, fusion-ready).

```
# Run the sweep (14 segments: 11 "big.bi" + 3 "bi.gbi")
python scripts/run_bigbi_polar_sweep.py

# Generate the dual-panel video (sagittal | polar)
python scripts/make_bigbi_dual_video.py

# Generate the self-contained HTML with embedded video
python scripts/embed_video_in_html.py en
```

The output is in output/bigbi_polar_sweep/:

- dual/dual_progressive.mp4 — dual-panel video (≈ 24.2 s, 1.4 MB)
- polar/polar_progressive.mp4 — polar-only video (3.3 MB)
- index_embedded.html — self-contained HTML (re-embedded in simulations/)

5.3 Demonstration 2: U-shaped T_p variation

The U-shaped variation lets the listener clearly perceive: **ib ib** at the start → **bi bi** at the bottom (fusion) → **ib ib** again at the end. This T_p sweep is the /ibi/ supplement demonstration (§3 classical transformation), where the inter-word pause IS the object of study; it is distinct from the big.bi sweep of Demonstration 1, which — in the dot form — carries no T_p dimension.

```
python scripts/run_ibbi_ushape_sweep.py
```

The sweep is a **continuous U-shaped ramp** of **10 segments**: T_p decreases to 0 with δ correlated at every point ($\delta = 0.5 + 0.5(1 - T_p/160)$), then increases back to the 160-ms baseline; consecutive segments flow directly into one another (no restarts):

	Segments	T_p (ms)	δ
descent	1–4	160, 120, 80, 40	0.50 ... 0.875
bottom (fusion, bi bi)	5	0	1.000
ascent	6–9	40, 80, 120, 160	0.875 ... 0.50
end baseline	10	160	0.500

Duration ≈ 18.6 s. Note: the coda @ is *not* held — the trajectory reaches the word-end anchor and decays immediately (no “ibe”).

5.4 Demonstration 3: Comparison with the article

To compare our polar figures with the article’s Figure 4:

```
# Generate the comparison figure
python scripts/compare_article_polar.py
```

This produces a 2×2 grid: article Figure 4 (top) vs our simulation (bottom), for both **big.bi** ($\delta = 1$) and **bi.gbi**.

6 Polar Trajectory Reconstruction

The two-branch polar visualization (vocalic z_v in red, consonantal z_c in blue) is rebuilt from the engine’s **recorded block sequence** by `scripts/polar_sync.py` (`record_pipeline` instruments the trajectory helpers without touching the engine files; `build_branches` then draws both branches from the blocks the engine actually walked). The original `Syllable_Synthesis` display anchoring is used — see `docs/DISPLAY_VS_ENGINE.md` for the full audit:

- every sub-arc carries the phase on the **arrival-side** point (p_a) with the departure-side point p_d held fixed (p_d = the consonant on z_c approach legs, V_1 on vocalic backgrounds) — this is the reference convention, not an option;
- the approach leg sweeps $\theta \in [0, \pi]$ ($\text{opint} = 0$), the release legs and vocalic backgrounds $\theta \in [-\pi, 0]$ ($\text{opint} = -1$, first sample dropped, for the releases);
- $\rho = \cos(\theta/2)$ ($P_{\text{exp}} = 1$).

6.1 Reconstruction rules

Block kind	Branch behavior
plateau, attack, decay	red: stationary anchor; blue: inactive
background, pause	red: arcplot arc ($V_1 \rightarrow V_2$) with $K = 30$, $\nu_v = -1$; blue: inactive
cluster ($n = (m+1)T$ steps)	red: arcplot arc ($V_1 \rightarrow V_2$) over the whole block; blue: sub-arcs $V_1 \rightarrow C_1 \rightarrow \dots \rightarrow C_m \rightarrow V_e$, T steps each, $K = 10$, $\nu_c = +1$, approach leg $[0, \pi]$ then release legs $[-\pi, 0]$ — the closed V→C→V gesture is a teardrop (belly at the vowel, point at the consonant), as in the article’s Figure 4

6.2 Display curvature and ν convention

The display uses the article values $K = 30$ for z_v and $K = 10$ for z_c (§2.1), and the per-branch display convention $\nu_v = -1$, $\nu_c = +1$ (the article’s planning-figure convention; the engine itself drives every arc with the single value $\nu = -1$). The legacy `polar_arc` form for z_v (mirrored phase sweep) remains available as `zv_form="polar_arc"` for the record.

6.3 Comparison with the article’s Figure 4

Under the article figure conditions ($\text{COEFCEN} = 1$, $\text{VOYDEB} = 0.5$, $T = 16$), the reproduced planning panels (`docs/figures/article_format/figure4_bigbi_bigbi.png`, `fig_bigbi_article_conditions.png` and `fig_words_article_conditions_nu_inv.png`) show:

- big.bi: **3** separate z_c gestures — /b/ (onset), palatal /g/ (coda), /b/ (word-2 onset, a closed loop anchored on the previous V_e);
- bi.gbi: **2** gestures — /b/ + the fused /gb/ cluster (single span visiting /g/ then /b/);
- the red z_v branch: schwa loop at the word onset (V_o at half radius), full vowel elsewhere ($V_e = V$ under $\text{COEFCEN} = 1$).

The article’s Figure 4 uses a different polar target table (the AP-model layout; see `docs/TRAJECTORY_MATCHING`), so the panel-level comparison is gestural (same teardrop topology, same gesture counts), not a point-wise overlay.

7 Results and Duration Analysis

7.1 Duration reduction

The article predicts a duration reduction when transitioning from **big.bi** ($\delta < 1$) to **bi.gbi** ($\delta = 1$):

“At first, the gain of consonant clustering is a reduction of one period T . The graph complexity is corollary reduced with a decay of the number of nodes with the loss of V .”

With $T = 16$, in the dot form (no pause — the block layout does not depend on δ , only the anchor values do), our simulation shows:

Configuration	δ	Steps	Duration (ms)	Reduction
big.bi (dot, start)	0.50	176	1760	—
big.bi (dot, pre-fusion)	0.995	176	1760	—
bi.gbi (fused)	1.00	160	1600	−160 ms (−1 T)

The observable fusion gain is $1T = 160$ ms, from consonant clustering: the fused /gb/ cluster chains the /g/→/b/ sub-arcs directly through the shared vowel anchor, and the separate syllable-2 onset gesture disappears. (The earlier $\approx 3T$ reading compared against the space form “big bi” included the inter-word pause machinery — decay + pause arc $\approx 2T$ — which, as established, is not part of big.bi: the ‘.’ is a C.C syllable boundary without pause, and the only latencies are the stop closures.)

7.2 Perceptual salience of the schwa @

With $T = 16$ and $\delta = 0.50$, the boundary schwa @ in **big.bi** is **salient** (clearly perceivable) because:

1. the C.C boundary anchor sits at $\rho = 0.45$ (most centralized, far from the full vowel /i/ at $\rho = 0.9$);
2. the z_v background dips to that anchor between the /g/ and /b/ clusters (audible schwa-like vocoid);
3. there is no pause: the schwa is carried by the planning trajectory through the consonant–consonant chain, exactly as the dot semantics prescribe.

As δ increases toward 1.0, the boundary anchor approaches $V = /i/$ (COEFCEN = 1) and the perception shifts from **big.bi** to **bi.gbi**.

8 Code Structure

8.1 Modified files (vs original timit-to-Maeda)

File	Modification
synthSYL/constants.py	Added DEFAULT_DELTA_0, DEFAULT_DELTA_E
synthSYL/gesture.py	Added δ_o , δ_e kwargs to build_gesture_anchors(), _insert_word_end_anchors(), _insert_vowel_onset_anchors(), _insert_syllable_onset_anchors() Fixed $V_o = (\delta_o \cdot \rho_V, \theta_V)$ (was $[\delta_o, \theta_V]$)
synthSYL/trajectory.py	Added V_e plateau (full-amplitude) before decay block
synthSYL/pipeline.py	Added δ_o , δ_e kwargs to panphon_pipeline()

8.2 Simulation scripts

Script	Purpose
run_bigbi_polar_sweep.py	Main sweep: big.bi → bi.gbi with polar coords
run_ibbi_ushape_sweep.py	U-shaped T_p variation
make_bigbi_dual_video.py	Dual-panel video generation
make_ibbi_polar_video.py	Polar-only video generation
make_article_figures.py	Article-format Figures 1/3/4 + article-conditions planning panels
embed_video_in_html.py	Self-contained HTML (base64 video)
compare_article_polar.py	Compare with article Figure 4

9 Quick Demo

The minimal code to run the big.bi → bi.gbi simulation:

```
from synthSYL import panphon_pipeline
import vlam
```

```

import numpy as np
from scipy.io.wavfile import write

# Run the pipeline for "big bi" at delta=0.5 (salient schwa)
result = panphon_pipeline(
    "big bi",
    T=16,
    short_pause_duration_factor=1.0,
    delta_o=0.5,
    delta_e=0.5,
)

# Synthesize audio
state = vlam.VlamState.initial(195)
config = vlam.SynthConfig(play_audio=False)
sr = vlam.synthwordfen(
    state=state,
    articulatory_params=result.Pval,
    word_tokens=["0", "v", "F"],
    duration_factor=16,
    envelope=result.envelope,
    config=config,
)

# Save WAV
sig = sr.signal
sig_int16 = np.int16(32767 * sig / (1.01 * np.max(np.abs(sig))))
write("big_bi_delta05.wav", 20000, sig_int16)
print(f"Signal: {sig.shape[0]} samples ({sig.shape[0]/20000*1000:.0f} ms)")

```

10 Article Demonstrations

This section documents three additional demonstrations from the article that have been implemented in this repository.

10.1 Demo 1: /ibia/ VCV synthesis (Figure 1)

The article’s Figure 1 shows the four-step synthesis pipeline on the sequence /ibia/ (/i/ → /b/ → /i/ → /a/). The four panels are:

1. **(a) Planning trajectories** $z_v(t)$ and $z_c(t)$ in the polar plane (ρ, θ) , article format. The vocalic branch z_v (red) flows from /i/ to /i/ to /a/, while the consonantal branch z_c (blue) makes a teardrop excursion toward the /b/ target.
2. **(b) Flow of articulatory parameters** — the 7 Maeda parameters over time. The *trough effect* is visible: the Body parameter (tongue body) makes an unexpected front/back movement during /b/, because /b/ lies on the same angular direction as /u/ in the complex plane (and coincides with it at the article position $\rho_b = 1$; see Demo 4).
3. **(c) Formant trajectory** F1-F2-F3. The F2 shows an S-shaped pattern: it starts at the /i/ value (≈ 1830 Hz), dips during /b/ (≈ 1564 Hz, the /b/ locus), then rises back to /i/ before transitioning to /a/.
4. **(d) Output spectrogram + wave** showing the acoustic result (waveform superposed on the grayscale spectrogram).

10.2 Demo 2: Locus equations (Figure 3)

The article’s Figure 3 shows the locus equations for /b/, /d/, /g/ across 8 vowels. A locus equation is the linear relationship:

$$F2_{\text{locus}} = a \cdot F2_{\text{vowel}} + b \quad (6)$$

where $F2_{\text{locus}}$ is measured 30 ms after the consonant **release** — the article: “taking the F2 values 30 ms after the onset” ($\delta_o = 0.5$), “onset” meaning the onset of the vowel at the release, in the locus-equation tradition of F2-onset measurement. In the synthesis timeline of a CV word ($T = 16$ steps of 10 ms), the release is the first frame of the C→V transition block (frame $2T$): frames $0..T-1$ are the initial arc (neutral → Vo), frames $T..2T-1$ the Vo→C closure (envelope silent), frames $2T..3T-1$ the C→V transition (envelope rising), and frames $3T..4T-1$ the vowel plateau. Hence $F2_{\text{locus}} = F2(\text{frame } 2T+3)$ and $F2_{\text{vowel}}$ is averaged over the plateau frames (envelope at full amplitude). The script (`scripts/run_locus_trough_demos.py`) detects the release and the plateau from the envelope; the offset (`LOCUS_OFFSET_STEPS = 3` in `config.py`) is the article’s 30 ms.

/g/ uses two article positions: velar ($\rho = 1.2, \theta = \pi/3$) for the back vowels {/u/, /o/, /O/} and palatal ($\rho = 1.1, \theta = 23\pi/12$) for the front vowels — *including* /a/, whose red point lies on the palatal branch in the article’s Figure 3.

Results (corrected protocol):

Consonant	Slope (a)	Intercept (b)	r^2	Article Fig. 3 (a)
/b/	0.704	351 Hz	0.982	≈ 0.70
/d/	0.412	1265 Hz	0.964	≈ 0.55
/g/ palatal (front + /a/)	0.644	903 Hz	0.809	≈ 0.75
/g/ velar (back)	2.107	−1207 Hz	0.999	≈ 2.0

The article-figure slopes in the last column were measured by pixel extraction from the arXiv source figure (marker centroids calibrated on the axis ticks). The /b/ slope matches the article to three digits, and the /g/ split (steep velar, shallow palatal) is reproduced. The /d/ locus at back vowels (1689–1743 Hz) sits in the canonical alveolar range (≈ 1700 –1800 Hz), and the slope ordering $a(/d/) < a(/b/) < 1$ now agrees with both the article’s Figure 3 and the classical locus-equation literature, where slopes descend from labial to velar to alveolar and the alveolar line is the flattest with the highest intercept [2, 3]. The article notes: “*This is similar to the classical observation but with a downward shift of the velar /g/ which could be due to geometric differences between the VLAM and the human vocal tract.*” The velar points here indeed fall on or below the $y = x$ diagonal (/g/+u: 1029 → 966 Hz).

The $F2_{\text{vowel}}$ axis is slightly compressed relative to the article’s figure (e.g. /i/: 2290 Hz here vs ≈ 2660 Hz, /u/: 1029 vs ≈ 880): the article was produced with an earlier synthSYL vowel table (corner vowels at $\rho = 1$ per §2.1), which has since been revised in `timit-to-Maeda`. This shifts the individual slopes slightly but preserves every trend.

10.2.1 Root cause of the previous /b/–/d/ inversion

An earlier version of this demo reported a spurious /b/–/d/ inversion (/d/ slope 1.633, /b/ 0.974). The cause was purely a **measurement-protocol bug**, not a model or synthesizer defect:

- The script derived the measurement frame from `anchor.i * T`, where `anchor.i` is the *node index* in the segment list (not a time). For a CV word the V anchor has $i = 1$, so F2 was sampled at frames $T+3$ (locus) and $T+8$ (“vowel”) — both *inside the silent Vo→C closure*, up to 13 frames before the release. The resulting regression of closure-phase F2 against closure-phase F2 produced the steep spurious /d/ line (reproduced exactly: $a = 1.633$, $b = -1316$, $r^2 = 0.991$).
- **The soft rectification is not implicated.** `valrect = 0.75` (`soft_rect_s`) is the canonical value of the original packages: `vlam.py` in this repository is byte-identical to `FBerthommier/timit-to-Maeda`, whose `batch_synthesize.py` defaults to `--valrect 0.75`. It is applied to the *area function* (`soft_rect` of section areas, preventing zero/negative areas during closure), not to the articulatory parameters. A sensitivity sweep confirms the locus slopes are insensitive to it over $s \in [0.5, 2.0]$ ($a(/d/) = 0.562 \pm 0.011$ at release+50 ms); only the hard limit $s \rightarrow 0$ (hard clipping) breaks the synthesis.
- δ_o **is not implicated either.** Sweeping $\delta_o \in \{0.5, 0.7, 1.0\}$ moves $a(/d/)$ by less than 0.02.
- **The measurement instant is the decisive parameter.** With the release correctly located at frame $2T$, the /d/ slope varies from 0.25 (release+0) to 0.79 (release+6 frames), passing through the article value around release+3 (30 ms). The /b/ slope varies from 0.50 to 0.92 over the same range.

No engine file (`vlam.py`, `synthSYL/`) was modified; only the measurement logic of the demo script, the /g/ allophony grouping of /a/, and `config.py` (new constant `LOCUS_OFFSET_STEPS`) changed.

Note: For a detailed comparison with Ohman’s spectrographic data and the S-shaped F2 trajectory analysis, see the 2025 companion article (arXiv:2505.22455).

10.3 Demo 4: Trough effect (Figure 1b)

The article (§3) describes an unexpected articulatory phenomenon during /b/:

“The tongue body is making an unexpected front/back movement during /b/ which corresponds to the trough effect [29].”

(reference [29]: Lindblom, Sussman, Modarresi Ghavami & Burlingame, *Phonetica* 59:245–262, 2002)

In the present model, /b/ lies on the **same angular direction as /u/** in the complex plane, with $\theta_b = \pi/3$. This is verified in the code constants: the vowel table of `synthSYL/phonology.py` places /u/ at ($\rho_u = 1.0$, $\theta_u = \pi/3$), and the article position of /b/ used by the demos is ($\rho_b = 1.0$, $\theta_b = \pi/3$) — with $\rho_b = \rho_u$ the two targets coincide exactly; the `synthSYL` registry default $\rho_b = 1.2$ shares the same direction but a larger radius. During the /b/ consonant:

- the consonantal branch z_c moves toward the /b/ target, which lies on the /u/ direction;
- the selection vector $S_c = \{1, 2, 6\}$ (article’s 1-based numbering; code indices $\{0, 1, 5\}$) activates the Jaw, Body and LipH parameters — in particular the Body parameter (index 1 in the 0-based code indexing);
- the vowel articulators selected by $S_v = \{3, 4, 5, 7\}$ (code indices $\{2, 3, 4, 6\}$: Dorsum, Tip, LipP, Hy) remain associated with the /i/ configuration;

- consequently, the Body parameter exhibits an unexpected front/back excursion during the /b/ interval, producing the characteristic dip in the simulated trajectory.

This articulatory excursion corresponds to the **trough effect** reported in Lindblom et al. [29]. In the present implementation, the effect arises from the interaction between the consonantal trajectory $z_c(t)$ and the parameter-selection mechanism: although /b/ is a bilabial consonant, its target lies in the same angular direction as /u/, and the Body parameter is included in the consonantal selection vector.

Thus, the simulation reproduces the tongue-body movement identified in the article as the trough effect. The geometric explanation given here is specific to the present model and should not be interpreted as the physiological explanation proposed by Lindblom et al. [29].

An interactive HTML simulation illustrating this effect accompanies this section in the `simulations/` directory.

Measured trough depth (for /ibi/, T=16):

- Body before /b/: ≈ 0.0 (neutral, at /i/)
- Body at trough: -2.25 (at $t = 160$ ms, during /b/)
- Body after /b/: ≈ 0.0 (back to /i/)
- Trough depth: 2.25 (significant dip)

This confirms the trough effect described in the article and in Lindblom et al. [29].

10.4 Demo 5: Six consonant cluster pairs

The article (§3) describes the tuning of the six pairs of /b,d,g/:

Pair	Input	Sc	Note
/bd/	abda	{1,2,3,6}	/b/ involved → Sc={1,2,3,6}
/bg/	abga	{1,2,3,6}	/b/ involved, /g/ palatal
/db/	adba	{1,2,3,6}	/b/ involved → Sc={1,2,3,6}
/dg/	adga	{1,2,3,4}	no /b/ → Sc={1,2,3,4}
/gb/	agba	{1,2,3,6}	/b/ involved, /g/ palatal
/gd/	agda	{1,2,3,4}	no /b/ → Sc={1,2,3,4}

Rule: when /b/ is involved in the pair, the selection vector is $S_c = \{1, 2, 3, 6\}$ (article numbering; code {0, 1, 2, 5} = jaw, body, dorsum, LipH); otherwise $S_c = \{1, 2, 3, 4\}$ (article numbering; code {0, 1, 2, 3} = jaw, body, dorsum, tip). The /g/ is systematically palatal when in a cluster.

10.5 Original article simulations (exact reproduction)

The `original_simulations/` directory contains three simulations reproduced with the **exact article parameters** for direct comparison with the article's figures:

10.5.1 Figure 1: /ibia/ with T=100ms

original_simulations/figure1_ibia/

Reproduces the article’s Figure 1 using $T = 10$ (100 ms per period), matching the article’s “ $T = 100\text{ms}$ ”. The 4-panel figure shows: (a) planning trajectories, (b) articulatory parameters (trough effect), (c) formant trajectory (S-shaped F2), (d) spectrogram.

10.5.2 Figure 4: big.bi and bi.gbi

original_simulations/figure4_bigbi_bigbi/ (raw script output) and docs/figures/article_format/ (article format).

big.bi is rendered with the **dot form** (big.bi): per the reference Timit-to-Maeda semantics (`_should_split_at_dot`: C.C → syllable split, C.V / V.C / V.V → merge), the dot between /g/ and /b/ is a syllable boundary with *no pause*, and the Ve of “big” is coarticulated with the /i/ of “bi” (syllable-onset anchor weighted by COEFCEN; full /i/ under COEFCEN = 1). Under the article figure conditions the left column thus shows three consonantal gestures (/b/; then /g/ and /b/ chained across the shared /i/) while the right column (bi.gbi, whose V.C dot is ignored) shows the fused /gb/ cluster — the §4 syllabification demonstration.

10.5.3 Supplement: /ib/ → /bi/

original_simulations/supplement_ib_bi/

Reproduces the supplement mentioned in §3: “The classical transformation of /ib/ into /bi/ is given in the supplement.” Three cases are shown side by side:

1. /ib/ with $\delta = 0.5$ (schwa @ audible, /b/ in coda)
2. /ib/ with $\delta = 1.0$ ($V_o = V_e = V$, fusion)
3. /bi/ with $\delta = 1.0$ (/b/ in onset, the fused form)

11 Interactive HTML Simulations

The repository includes three self-contained HTML files (in the `simulations/` directory) with video, poster, captions, and comparison figures embedded as base64 data URIs. These files can be opened directly in any modern browser without requiring a web server or external file dependencies.

11.1 big.bi → bi.gbi (T=16, starts at COEFCEN=0.5)

simulations/bigbi_bi_gbi_embedded.html (≈ 6.3 MB)

This simulation demonstrates the verbal transformation from big.bi to bi.gbi (§4 of the article), in the **dot form** (no pause, no T_p dimension). It starts at $\delta = 0.50$ (COEFCEN = 0.5, salient boundary schwa @) then increases δ to 1.00 (the C.C boundary anchor reaches the full /i/, fusion). The last bi.gbi segment is repeated 3 times for perceptibility. Duration ≈ 24.2 s.

Features: dual-panel video (sagittal Maeda | polar plot with phoneme), 14 segments, WebVTTCaptions, indexed transcript, keyboard shortcuts.

11.2 ib ib → bi bi (NON-REVERSIBLE U-shape)

`simulations/ibbi_ushape_nonreversible_embedded.html` (≈ 4.4 MB)

This simulation demonstrates the **non-reversible** verbal transformation with the **same U-shaped T_p as the reversible demonstration**: “ib ib” (VC.VC, /b/ in coda) with T_p decreasing 160, 120, 80, 40, 10 ms (δ correlated at every point); at the fusion point ($T_p = 0$, $\delta = 1$) the input switches to “bi bi” (CV.CV, /b/ in onset); then T_p increases again 40, 80, 120, 160 ms — but the content STAYS “bi bi”: the /b/ remains in onset position even with a fully restored pause, proving the transformation is permanent and non-reversible. Duration ≈ 16.9 s (10 segments).

Features: dual-panel video, 10 segments, WebVTT captions, indexed transcript, keyboard short-cuts.

11.3 ib ib → bi bi (U-shape, reversible)

`simulations/ibbi_ushape_embedded.html` (≈ 5.3 MB)

The reversible version: a **continuous U-shaped ramp** of 10 segments (see Demonstration 2) — “ib ib” throughout, with T_p 160 → 0 → 160 ms, δ correlated. This version returns to “ib ib” at the end (reversible). Included for comparison with the non-reversible version above. Duration ≈ 18.6 s.

Features: dual-panel video, 10 segments, WebVTT captions, indexed transcript, keyboard short-cuts.

12 Conclusion

This repository provides a complete, reproducible implementation of the verbal transformation described in arXiv:2307.02299. The key contributions are:

1. The δ_o and δ_e parameters, enabling the gradual transition from schwa-like vocoid to full vowel.
2. The syllable-onset anchor weighted by COEFCEN (the reference Timit-to-Maeda formula), and no held @: the trajectory reaches the word-end anchor and decays immediately (one does not pronounce “ibe”).
3. The V_o formula correction ($V_o = \delta_o \cdot \rho_V$), ensuring $V_o = V$ at $\delta = 1$ for proper fusion.
4. The polar trajectory reconstruction, reproducing the two-branch visualization (vocalic z_v , consonantal z_c) described in the covtl-pipeline documentation.
5. The U-shaped T_p variation, for clear perceptual demonstration of the ib ib → bi bi transformation.

The duration reduction observed ($3.06T \approx 490$ ms with $T = 16$) matches the article’s prediction of $3T$ ($1T$ from consonant clustering + $2T$ from V_o/V_e plateau absorption).

The polar topology matches the article’s Figure 4: at $\delta < 1$, the consonantal branch makes 2 separate excursions; at $\delta = 1$ (bi.gbi), it makes 1 fused /gb/ excursion.

References

- [1] F. Berthommier, *Why can big.bi be changed to bi.gbi? A mathematical model of syllabification and articulatory synthesis*, arXiv:2307.02299, July 2023.
- [2] H. M. Sussman, H. A. McCaffrey and S. A. Matthews, *An investigation of locus equations as a source of relational invariance for stop place categorization*, J. Acoust. Soc. Am. 90(3):1309–1325, 1991.
- [3] K. Iskarous, C. A. Fowler and D. H. Whalen, *Locus equations are an acoustic expression of articulator synergy*, J. Acoust. Soc. Am. 128(4):2021–2032, 2010.

Figure 1: four-step synthesis of /ibia/ — $T = 100$ ms ($K=10$, $K_{\text{voy}}=30$, $P_{\text{exp}}=1$)
 arXiv:2307.02299 (Berthommier 2023)




Figure 1: Figure 1, article format (`make_article_figures.py`): four-step synthesis of /ibia/ with $T = 100$ ms ($T_{\text{base}} = 10$, $K=10$, $K_{\text{voy}}=30$, $P_{\text{exp}}=1$). (a) Polar planning trajectories (z_v red, z_c blue — original arcplot anchoring, teardrop z_c). (b) The 7 Maeda parameters with phoneme ticks /i/ /b/ /i/ /a/ (trough effect visible on Body). (c) Formant trajectory (S-shaped F2). (d) Spectrogram + wave. Compare with the article's Figure 1 (side-by-side sheet: `P:\bigbi-workspace\sheet_fig1.jpg`, kept outside the repository).




Figure 2: Figure 3, article format (make_article_figures.py): locus equations for /b/, /d/, /g/ across 8 vowels. Each panel shows $F2_{\text{locus}}$ vs $F2_{\text{vowel}}$ with the linear regression line.




Figure 3: Figure 3 (combined), article format: locus equations for all three consonants. The dotted line ($y = x$) represents no coarticulation. Slopes < 1 indicate coarticulation (the consonant locus is influenced by the vowel context).




Figure 4: Trough effect: Body parameter variation during /b/, /d/, /g/ in VCV sequences. The Body (orange, thick line) shows a dip during the consonant for /b/ (ibi, aba), corresponding to the trough effect. The consonant region is shaded in orange.




Figure 5: Body parameter overlay for /ibi/, /aba/, /idi/, /igi/. The trough (dip) is most pronounced for /b/ (ibi, aba) because the /b/ target lies on the same angular direction as /u/.




Figure 6: Polar explanation: /b/ (orange square) lies on the same angular direction as /u/ (orange circle) in the complex plane — with the article position $\rho_b = 1$ the two targets coincide. During /b/, the Body parameter follows $z_c(t)$ toward /u/, creating the trough effect while the vowel articulators stay around /i/.




Figure 7: Articulatory parameters for the six cluster pairs. Each panel shows the 7 Maeda parameters over time.




Figure 8: Polar trajectories for the six cluster pairs. The consonantal branch z_c (blue) visits both consonant targets in sequence.




Figure 9: Spectrograms for the six cluster pairs.




Figure 10: Figure 1 (original reproduction): four-step synthesis of /ibia/ with $T = 100\text{ms}$ ($T_{\text{base}} = 10$).




Figure 11: Figure 4, article format (make_article_figures.py): big.bi | bi.gbi under the article figure conditions $\text{COEFCEN} = 1$ (δ_e), $\text{VOYDEB} = 0.5$ (δ_o), $T = 16$, both in the **dot form** (no pause). big.bi is the C.C syllable concatenation: the V_e of “big” is coarticulated with the $/i/$ of “bi” (the syllable-onset anchor is the full previous vowel under $\text{COEFCEN} = 1$) — the only silences are the stop closures ($\approx T/2$), as in the reference realization; z_c carries 3 consonantal gestures ($/b/$; $/g/$ then $/b/$ chained across the shared $/i/$). bi.gbi: the dot between $/i/$ and $/g/$ is V.C and therefore ignored (syllables merged), giving the fused $/gb/$ cluster (2 gestures). Top: polar planning panels (teardrop z_c , red z_v holding the $/i/$ radius). Middle: Maeda parameters with phoneme ticks. Bottom: wave (ms) + spectrogram (s), same total duration. Compare with the article’s Figure 4 (side-by-side sheet kept outside the repository: P:\bigbi-workspace\sheet_fig4.jpg).




Figure 12: Supplement (original reproduction): /ib/ → /bi/ classical transformation. Three columns: /ib/ $\delta=0.5$, /ib/ $\delta=1.0$, /bi/ $\delta=1.0$.